

USI Dynamic Analysis Group Project

Dynamically executed SIMD instructions profiler

Introduction (1/2) - Project

Statement

- Count the **executed** machine vector instructions (**SIMD** instructions) during a certain *Java Virtual Machine (JVM)* execution

Requirements

- ✓ The solution should **report** the **numbers** of executed vector instructions using **counters**
- ✓ The solution must **work** on *Linux Ubuntu*
- ✓ The solution must **work** on *OpenJDK 18*

Introduction (2/2) - Features

Not-Available features

- ✗ Consider **only** the machine vector instructions **emitted** by the *Just-In-Time (JIT) compiler* (i.e. **C2**) (partially obtained, requires a more in-depth analysis of the JVM architecture)
- ✗ **Expose** a *Java API* to **reset** counters

Available features

- ✓ Report **different** counters for different **types** of instructions (e.g. *load, store, arithmetics, etc.*)
- ✓ Report **thread-local** counters (instead of a *single global counter per type*)

Intel Pin (1/1) - Analysis

- **Platform** for creating **analysis tools** that comprises the following types of routines
 - **Instrumentation** - are called when code is **about** to be run (**before** recompilation) and **enable** the insertion of Analysis routines
 - **Analysis** - are called when the code **associated** with them is run
 - **Callback** - are called when specific **conditions** are met or when a certain **event** has occurred
- **Portable** architecture by isolating platform-specific code from generic code
- Performs a large collection of **optimization** techniques to **reduce** running time and memory overhead
- **Used** in **JIT mode** (it uses a **JIT compiler** to recompile the code and insert instrumentation)

Pintool (1/9) - Headers (Portability)

- A local copy of **Intel Pin 3.24 98612** is **required**
- **pin.H** is the **entry-point** header of the entire library
- The rest of the `#includes` refer to the **portable CXX stdlib** headers **exposed** by the library

```
// PIN_ROOT/source/include/pin/pin.H
#include "pin.H"

// PIN_ROOT/extras/cxx/include/...
#include "fstream"
#include "iostream"
#include "unordered_map"
```

Pintool (2/9) - Data structures (x86 Extension sets)

- Describes the **available x86 Extension Sets**
- More sets can be **added** here
- Used to **filter** the **instructions** by the required sets
- The **_NONE_** variant is used to describe an instruction that does **not belongs** to the defined Extension sets

```
/// Type: x86 Extension Sets
/// @note Add more sets here
enum class VXESet {
    _NONE_ = -1, ///< Not an extension
    SSE1f,      ///< SSE (Pentium 3, 1999), Floating-point
    SSE1i,      ///< SSE (Pentium 3, 1999), Integer
    SSE2f,      ///< SSE2 (Pentium 4), Floating-point
    SSE2i,      ///< SSE2 (Pentium 4), Integer
    SSE3,       ///< SSE3 (later Pentium 4)
    SSSE3,      ///< SSSE3 (early Core 2)
    SSE41,      ///< SSE4.1 (later Core 2)
    SSE4a,      ///< SSE4.a (Phenom)
    SSE42,      ///< SSE4.2 (Nehalem)
    MMX,        ///< MMX (1996)
};
```

Pintool (3/9) - Data structures (x86 Instruction types)

- Describes the **available x86 Instruction types**
- More types can be **added** here
- Used to **filter** the **instructions** by type
- The **_NONE_** variant describes an **unknown** type (should not be used)

```
/// Type: x86 Instruction Types
/// @note Add more types here
enum class VXType {
    _NONE_,
    LOAD,
    STORE,
    DATA_TRANSFER,
    CONVERSION,
    ARITHMETIC,
    COMPARISON,
    LOGICAL,
    EXTRACT,
    INSERT,
    SHUFFLE,
    SHIFT,
    PACK,
    UNPACK,
    STATE_MANAGEMENT,
};
```

Pintool (4/9) - Data structures (Associations)

- Definition of the **associations** between **Instructions** and **Extension sets, Descriptions and Types**
- The **keys** are **Instruction Opcodes** `UINT16` from `typedef enum { ... } xed_iclass_enum_t` contained in `$PIN_ROOT/extras/xed-intel64/include/xed/xed-iclass-enum.h`
- More associations can be **added** here
- The *Intel Pin* **API** does **not** natively offer this fine-grained filtering
- **Improvements** - **perfect hashing** can be implemented

```
/// Table: Instruction -> Extension Set
/// @note Add more entries here
const std::unordered_map<UINT16, VXSet> VXTableESet = {
    // SSE (Pentium 3, 1999), Floating-point
    { XED_ICCLASS_ADDSS,    VXSet::SSE1f },
    { XED_ICCLASS_ADDPS,    VXSet::SSE1f },
    { XED_ICCLASS_CMPPS,    VXSet::SSE1f },
```

```
/// Table: Instruction -> Description
/// @note Add more entries here
const std::unordered_map<UINT16, std::string> VXTableDesc = {
    // SSE (Pentium 3, 1999), Floating-Point
    { XED_ICCLASS_ADDSS,    "Add Scalar Single-Precision Float"},
    { XED_ICCLASS_ADDPS,    "Add Packed Single-Precision Float"},
    { XED_ICCLASS_CMPPS,    "Compare Packed Single-Precision F
```

```
/// Table: Instruction -> Type
/// @note Add more entries here
const std::unordered_map<UINT16, VXType> VXTableType = {
    // SSE (Pentium 3, 1999), Floating-Point
    { XED_ICCLASS_ADDSS,    VXType::ARITHMETIC },
    { XED_ICCLASS_ADDPS,    VXType::ARITHMETIC },
    { XED_ICCLASS_CMPPS,    VXType::COMPARISON },
```


Pintool (5/9) - Data structures (Thread local data)

- Definition of the **thread-local data**
- The **counters** are defined inside a **unordered_map** (**Extension sets** driven) of **unordered_map** (**Instruction opcodes** driven)
- The data cache lines are **separated** to **avoid** race conditions (multithreading)
- The data are **retrieved** through a **TLS storage key** (initialized once in **main**)

```
// Alias definition -> umap (extension-set: (opcode: counter))
using VXTableCount =
    std::unordered_map<VXSet, std::unordered_map<UINT16, UINT64>>>;

/// Thread's data = id + counters
/// Let each thread's data be in its own data cache line so that
/// multiple threads do not contend for the same data cache line
struct ThreadData {
    THREADID id;
    VXTableCount tc;
    ThreadData();
};
```

```
// Key for accessing TLS storage in the threads
// @note Initialized once in main()
static TLS_KEY tls_key;

/// Function to access thread-specific data
/// @param tid current thread id (assigned by pin)
ThreadData *getTLS(THREADID tid) {
    return static_cast<ThreadData *>(PIN_GetThreadData(tls_key, tid));
}
```

Pintool (6/9) - Callback routines (Thread related)

- The `threadStart` hook is called for every thread created by the application (when it is about to start) and performs the `ThreadData` construction and initialization
- The `threadFini` hook is called for every thread destroyed by the application and prints out `ThreadData` in **CSV** format

```
/// This function is called for every thread destroyed
/// by the application
/// Print out analysis results:
/// - The data of threads to:
///   - <name>.td.csv OR
///   - stdout
/// @param tid thread id (assigned by pin)
/// @param ctxt initial register state for the new thread
/// @param flags thread creation flags (OS specific)
/// @param v value specified by the tool in the
///       PIN_AddThreadFiniFunction call
VOID threadFini(THREADID tid, const CONTEXT *ctxt, INT32 code, VOID *v) {
    *td_csvOut << *getTLS(tid);
}
```

```
/// This function is called for every thread created
/// by the application when it is about to start
/// running (including the root thread)
/// @param tid thread id (assigned by pin)
/// @param ctxt initial register state for the new thread
/// @param flags thread creation flags (OS specific)
/// @param v value specified by the tool in the
///       PIN_AddThreadStartFunction call
VOID threadStart(THREADID tid, CONTEXT *ctxt, INT32 flags, VOID *v) {
    // Increase the number of threads counter
    NThreads++;

    // abort() if NThreads > maxNThreads
    // could be an ASSERT() call
    if (NThreads > maxNThreads) {
        std::cerr << "max number of threads exceeded!" << std::endl;
        PIN_ExitProcess(1);
    }

    // Create new ThreadData
    ThreadData *data = new ThreadData();
    PIN_SetThreadData(tls_key, data, tid);
    // Assign id
    data->id = tid;
}
```

Pintool (7/9) - Instrumentation routine

- The **trace** function is called every time a new **trace** is encountered
- The function **inserts** a **call** to the **VXCountIncr** *analysis routine*
- **IPOINT_ANYWHERE** is used to let *Pin* **schedule** the call anywhere and obtain best **performance**
- The function visit every **basic block** in the trace and every **instruction** inside it
- The function inserts the *analysis routine* call **only** if the current instruction **belongs** to one of the defined **Extension sets** (**filtering**)

```
/// This function is called every time a new trace is encountered
/// It inserts a call to the VXCountIncr analysis routine
/// @param trace trace to be instrumented
/// @param value specified by the tool in the
/// TRACE_AddInstrumentFunction call
VOID trace(TRACE trace, VOID *v) {
    // Visit every basic block in the trace
    for (BBL bbl = TRACE_BblHead(trace); BBL_Valid(bbl); bbl = BBL_Next(bbl)) {
        // Visit every instruction in the current basic block
        for (INS ins = BBL_InsHead(bbl); INS_Valid(ins); ins = INS_Next(ins)) {
            // Get the current instruction opcode
            OPCODE insOpcode = INS_Opcode(ins);
            // Get the current instruction extension-set
            VXESet insESET = getVXESet(insOpcode);
            // If the current instruction is a vector instruction
            if (insESET != VXESet::_NONE_)
                // Insert a call to VXCountIncr passing the opcode
                // and the set of the instruction
                // IPOINT_ANYWHERE allows Pin to schedule the call
                // anywhere to obtain best performance
                BBL_InsertCall(bbl, IPOINT_ANYWHERE, (AFUNPTR)VXCountIncr,
                             IARG_FAST_ANALYSIS_CALL, IARG_UINT32, insOpcode,
                             IARG_UINT32, insESET, IARG_THREAD_ID, IARG_END);
        }
    }
}
```

Pintool (8/9) - Analysis routine

- The `VXCountIncr` function is called for every `filtered instruction` within each `basic block`
- The function `increments` by 1 the correct `thread-local counter`
- The `macro` `PIN_FAST_ANALYSIS_CALL` is used to let *Pin* perform a `faster linkage` for this call, increasing performances

```
/// This function is called for every basic block
/// Increments the correct thread-local counter:
/// - extension-set -> opcode -> counter
/// @param o current instruction opcode
/// @param x current instruction extension-set
/// @param x current thread id (assigned by pin)
/// @note use atomic operations for multi-threaded applications
VOID PIN_FAST_ANALYSIS_CALL VXCountIncr(OPCODE opcode, VXESet eset,
                                         THREADID tid) {
    getTLS(tid)->tc[eset][opcode]++;
}
```

Pintool (9/9) - Application lifecycle

- Defined within the `main` function (not shown entirely)
- Registration of the `threadStart` and `threadFini` hooks to be called at the creation and destruction of a thread, respectively
- Registration of the `fini` hook (not showed) to be called when the application exits (simply prints out the number of created threads in **CSV** format)
- Registration of the `trace` hook to be called to instrument filtered instructions
- Call to the `divergent` function `PIN_StartProgram`, that `never` returns

```
// Register threadStart to be called when thread starts
PIN_AddThreadStartFunction(threadStart, nullptr);
// Register threadFini to be called when thread exits
PIN_AddThreadFiniFunction(threadFini, nullptr);
// Register fini to be called when the application exits
PIN_AddFiniFunction(fini, nullptr);
// Register trace to be called to instrument instructions
TRACE_AddInstrumentFunction(trace, nullptr);

// Divergent function: never returns
PIN_StartProgram();
```

Usage (1/2) - Infrastructure

Description

- Simple make driven **infrastructure** to test the solution (all the following details are automatically handled within the Makefile)
- Few **dependencies** with standard tools (e.g. make, coreutils, linux-utils, etc.)
- To compile the **Pintool** a C++ **compiler**, a local **copy** of *Intel Pin 3.24 98612* and the definition of the `PIN_ROOT` *environment variable* are **required**
- **JVBench** (*JVBench-1.0.jar*) is used for **testing** and all the tests in the suite are **available** (*OpenJDK >= 16* is **required**)

Testing

- Simply **call** make followed by a (list) target test **name(s)** (call make or make help for info)
- The `jdk.incubator.vector` module is automatically included
- A simple Python **script** for **bulk** executions is also provided

Usage (2/2) - Testing (example)

- Automatically **builds** the **dependencies** (*Intel Pin* and *JVBench*) **environment** within the `env` folder
- **Compiles** the Pintool
- **Instrument** and **run** the target test application (the execution is **silenced** and the result is redirected to a **log** file)
- **Prints** out the test results (**CSV**)
- **Moves** the test results within the `test` folder

```
dxvc # make AxyBenchmark.autoVec
Creating Directory /home/resu/workspace/github/resu-gh/dxvc/env
Downloading Intel Pin to /home/resu/workspace/github/resu-gh/dxvc/env/intelpin
Unzipping /home/resu/workspace/github/resu-gh/dxvc/env/intelpin.tar.gz
Cloning uci-dag/JVBench to /home/resu/workspace/github/resu-gh/dxvc/env/jvbench
Downloading JVBench-1.0.jar to /home/resu/workspace/github/resu-gh/dxvc/env/jvbench
Compiling... /home/resu/workspace/github/resu-gh/dxvc/src/dxvc/dxvc.cpp
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
egrep: warning: egrep is obsolete; using grep -E
Running benchmark AxyBenchmark.autoVec
INFO: This application is instrumented by dxvc
INFO: Writing to: AxyBenchmark_autoVec.[td,nt].csv
WARNING: Using incubator modules: jdk.incubator.vector
Creating Directory /home/resu/workspace/github/resu-gh/dxvc/test/AxyBenchmark_autoVec
Moving AxyBenchmark_autoVec.[td,nt].csv to /AxyBenchmark_autoVec
Displaying Results ...
thread set opcode mnemonic category extension type counter description
17 SSE2F 475 MULSD --- --- ARITHMETIC 45 Multiply Scalar Double-Precision Floating-Point Values
17 SSE2F 124 COMISD --- --- COMPARISON 29 Compare Scalar Ordered Double-Precision Floating-Point Values
17 SSE2F 428 MOVAPD --- --- DATA_TRANSFER 11 Move Aligned Packed Double-Precision Floating-Point Values
17 SSE2F 417 MAXSD --- --- COMPARISON 10 Return Maximum Scalar Double-Precision Floating-Point Value
17 SSE2F 160 DIVSD --- --- ARITHMETIC 35 Divide Scalar Double-Precision Floating-Point Values
17 SSE2F 149 CVTTS2SDSI --- --- CONVERSION 10 Convert with Truncation Scalar Double-Precision Floating-Point Values
17 SSE2F 141 CVTSD2SD --- --- CONVERSION 98 Convert Doubleword Integer to Scalar Doubleword Integer
17 SSE2F 858 UCOPUSD --- --- COMPARISON 20 Unordered Compare Scalar Double-Precision Floating-Point Values
17 SSE2F 11 ADDSD AVX512_BITALG AVX2GATHER ARITHMETIC 1 Add Scalar Double-Precision Floating-Point Values
17 SSE2F 143 CVTSS2SD --- --- CONVERSION 9 Convert Scalar Single-Precision Floating-Point Values to Double-Precision Floating-Point Values
17 SSE2I 665 BUNPQKHQQD --- --- UNPACK 2 Unpack High Data
17 SSE2I 510 PADDD --- --- ADD_PCKED_INTEGER 8 Add Packed Integers
17 SSE2I 437 MOVDQB --- --- DATA_TRANSFER 1551 Move Unaligned Packed Integer Values
17 SSE2I 634 PSHUFD --- --- SHUFFLE 128 Shuffle Packed Doublewords
threads
49
dxvc #
```


Conclusion

Advantages

- The Pintool is **not** strictly **binded** with the **JVM** (e.g. execution on a simple “hello world” Python script)

Limitations

- **JVBench** tests CLI **arguments** are not dynamically managed (use an environment variable)
- java CLI **arguments** are not dynamically managed (use an environment variable)
- Test results are not **discriminated** (add timestamps to filenames)
- **Portability** could be improved (e.g. Docker, nix, etc.)

```
dxvc # bat --plain main.py
print('hello world')
dxvc # ./env/intelpin/pin -t src/dxvc/obj-intel64/dxvc.so -- python main.py
INFO: This application is instrumented by dxvc
INFO: Writing to: stdout
thread,set,opcode,mnemonic,category,extension,type,counter,description
hello world
0,SSE41,748,ROUNDSD,---,---,ARITHMETIC,1,Round Scalar Double Precision Floatin
0,SSE4a,412,LZCNT,---,---,COMPARISON,44,Count the Number of Leading Zero Bits
0,SSE42,538,PCMPISTRI,---,---,COMPARISON,3,Packed Compare Implicit Length Stri
0,SSSE3,516,PALIGNR,---,---,ARITHMETIC,23,Packed Align Right
0,SSSE3,633,PSHUFB,---,---,SHUFFLE,2,Packed Shuffle Bytes
0,SSE2f,858,UCOMISD,---,---,COMPARISON,30,Unordered Compare Scalar Double-Prec
AGS
```


Thanks for your attention